

Underactuated Robots

Lecture 3: Differential Dynamic Programming

Nicola Scianca

September 2024

- the Bellman equation is at the core of **dynamic programming**, a technique for solving optimization problems by breaking them down into smaller subproblems
- it defines a **recursive relationship**, where the expected value at a particular time depends on the value at a future time
- it has important applications not only in optimal control and trajectory optimization, but also in reinforcement learning and other areas

the value function

$$V_k(x_k) = \min_{u_k, \dots, u_{N-1}} \left(\sum_{i=k}^{N-1} l_k(x_i, u_i) + l_N(x_N) \right)$$

s. t. $x_{k+1} = f(x_k, u_k)$

- the **value function** $V_k(x_k)$ tells us the cost we will pay if we always take optimal actions in the future
- the cost is a reward in maximization problems, hence the name “value function”
- at time k , it depends on the state x_k , but not on the input (we are choosing the best possible input sequence)

the Bellman equation

- now, take out the first item of the sum

$$\begin{aligned} V_k(x_k) &= \min_{u_k, \dots, u_{N-1}} \left(l_t(x_t, u_t) + \sum_{i=k+1}^{N-1} l_k(x_i, u_i) + l_N(x_N) \right) \\ &= \min_{u_k} \left(l_t(x_t, u_t) + \min_{u_{t+1}, \dots, u_{N-1}} \sum_{i=k+1}^{N-1} l_k(x_i, u_i) + l_N(x_N) \right) \\ &= \min_{u_k} \left(l_t(x_t, u_t) + V_{k+1}(x_{k+1}) \right) \end{aligned}$$

- since V_{k+1} is the value function at $k+1$, we get a nice recursive expression: this is the Bellman equation

the Bellman equation

- remember: this minimization must take into account the dynamic constraint $x_{k+1} = f(x_k, u_k)$
- if we substitute this inside the recursive expression, we get

$$V(x_k) = \min_{u_k} \left(l_t(x_t, u_t) + V_{k+1}(f(x_k, u_k)) \right)$$

which is the Bellman equation in the discrete case

the Bellman equation

$$V(x_k) = \min_{u_k} \left(l_t(x_t, u_t) + V_{k+1}(f(x_k, u_k)) \right)$$

- the Bellman equation relates the value of V_k to the value of V_{k+1} in a recursive way
- if we had perfect knowledge of V , this would make it easy to derive an optimal control law, but usually we cannot explicitly solve for V and we have to approximate
- however, there is at least one case in which we can find an explicit solution: the **Linear Quadratic Regulator** (LQR)

linear quadratic regulator

- in the LQR we are trying to regulate the state x to zero, thus the cost function to minimize looks like

$$J = \sum_{k=0}^{N-1} (x_k^T Q x_k + u_k^T R u_k) + x_N^T Q_f x_N$$

where $Q \succeq 0$, $R \succ 0$, and $Q_f \succeq 0$ are weight matrices

- the system dynamics equation is linear

$$x_{k+1} = Ax_k + Bu_k$$

linear quadratic regulator

- the value function at time step k is

$$V_k(x_k) = \min_{u_k, \dots, u_{N-1}} \sum_{i=k}^{N-1} (x_i^T Q x_i + u_i^T R u_i) + x_N^T Q_f x_N$$

- using the Bellman equation we can write it recursively as

$$V_k(x_k) = \min_{u_k} (x_k^T Q x_k + u_k^T R u_k + V_{k+1}(x_{k+1}))$$

with system dynamics $x_{k+1} = Ax_k + Bu_k$.

linear quadratic regulator

- let's assume that the value function is **quadratic**

$$V_k(x_k) = x_k^T P_k x_k$$

- the value function at the next time-step $k + 1$ can be related to the current x_k and u_k via the system dynamics

$$V_{k+1} = x_{k+1}^T P_{k+1} x_{k+1} = (Ax_k + Bu_k)^T P_{k+1} (Ax_k + Bu_k)$$

- substitute into the Bellman equation:

$$V_k(x_k) = \min_{u_k} \left(x_k^T Q x_k + u_k^T R u_k + (Ax_k + Bu_k)^T P_{k+1} (Ax_k + Bu_k) \right)$$

linear quadratic regulator

- to compute the minimum, derive with respect to u_k and set equal to zero

$$\begin{aligned} & \frac{d}{du_k} \left(x_k^T Q x_k + u_k^T R u_k + (A x_k + B u_k)^T P_{k+1} (A x_k + B u_k) \right) \\ &= 2 R u_k + 2 B^T P_{k+1} (A x_k + B u_k) \\ &= (R + B^T P_{k+1} B) u_k + B^T P_{k+1} A x_k = 0 \\ &\implies u_k = -(R + B^T P_{k+1} B)^{-1} B^T P_{k+1} A x_k \end{aligned}$$

- substituting into the Bellman equation yields the **Riccati recursion**

$$P_k = Q + A^T P_{k+1} A - A^T P_{k+1} B (R + B^T P_{k+1} B)^{-1} B^T P_{k+1} A$$

linear quadratic regulator

- now we start from the final state with $P_N = Q_f$, and going backwards with the Riccati recursion we can compute P_k at each time step

$$P_k = Q + A^T P_{k+1} A - A^T P_{k+1} B (R + B^T P_{k+1} B)^{-1} B^T P_{k+1} A$$

- once we have full knowledge of P_k (which means full knowledge of the value function), the optimal control law is simply given by

$$u_k = -(R + B^T P_{k+1} B)^{-1} B^T P_{k+1} A x_k$$

- if the systems dynamics are nonlinear, we usually can't directly solve for the value function
- if we discretize the state space we can compute V numerically everywhere (dynamic programming), but for large systems this is impossible due to the **curse of dimensionality**
- what we can do is find a **quadratic approximation** of the value function: **Differential Dynamic Programming** (DDP)

quadratic approximation of the value function

- let's recall the Bellman equation

$$V(x_k) = \min_{u_k} \left(l_t(x_t, u_t) + V_{k+1}(f(x_k, u_k)) \right)$$

- call Q the argument of the minimization

$$Q(x_k, u_k) = l_t(x_t, u_t) + V_{k+1}(f(x_k, u_k))$$

quadratic approximation of the value function

- a quadratic approximation Q , around the point (\bar{x}_k, \bar{u}_k) , would look like

$$\begin{aligned} Q(x_k, u_k) \simeq & Q(\bar{x}_k, \bar{u}_k) + \left. \frac{dQ}{dx_k} \right|_{\bar{x}_k, \bar{u}_k} \Delta x_k + \left. \frac{dQ}{du_k} \right|_{\bar{x}_k, \bar{u}_k} \Delta u_k \\ & + \frac{1}{2} \Delta x_k^T \left. \frac{d^2 Q}{dx_k^2} \right|_{\bar{x}_k, \bar{u}_k} \Delta x_k + \frac{1}{2} \Delta u_k^T \left. \frac{d^2 Q}{du_k^2} \right|_{\bar{x}_k, \bar{u}_k} \Delta u_k \\ & + \Delta u_k^T \left. \frac{d^2 Q}{dx_k du_k} \right|_{\bar{x}_k, \bar{u}_k} \Delta x_k \end{aligned}$$

- let's use a more compact notation

$$\begin{aligned} \tilde{Q}(x_k, u_k) = & \bar{Q}_k + \bar{Q}_k^x \Delta x_k + \bar{Q}_k^u \Delta u_k \\ & + \frac{1}{2} \Delta x_k^T \bar{Q}_k^{xx} \Delta x_k + \frac{1}{2} \Delta u_k^T \bar{Q}_k^{uu} \Delta u_k + \Delta u_k^T \bar{Q}_k^{ux} \Delta x_k \end{aligned}$$

quadratic approximation of the value function

- let's compute, as an example, \bar{Q}_k^x

$$\bar{Q}_k^x = \frac{\partial}{\partial x_k} \left(l_k(x_k, u_k) + V_{k+1}(f(x_k, u_k)) \right) = \frac{\partial l}{\partial x_k} \Big|_{\bar{x}_k, \bar{u}_k} + \frac{\partial V_{k+1}}{\partial x_k} \Big|_{\bar{x}_k, \bar{u}_k}$$

- since V_{k+1} is the value function at $k+1$, it doesn't depend on x_k directly but through the dynamics $f(x_k, u_k)$; therefore, we must use the chain rule

$$\bar{Q}_k^x = \frac{\partial l}{\partial x_k} \Big|_{\bar{x}_k, \bar{u}_k} + \frac{\partial V_{k+1}}{\partial x_{k+1}} \Big|_{\bar{x}_k, \bar{u}_k} \frac{\partial f}{\partial x_k} \Big|_{\bar{x}_k, \bar{u}_k} = \bar{l}_k^x + V_{k+1}^x \bar{f}_k^x$$

quadratic approximation of the value function

- similarly we can compute all the other terms

$$\bar{Q}_k^x = l_k^x + V_{k+1}^x f_k^x$$

$$\bar{Q}_k^u = l_k^u + V_{k+1}^x f_k^u$$

$$\bar{Q}_k^{xx} = l_k^{xx} + (f_k^x)^T V_{k+1}^{xx} f_k^x + \cancel{V_{k+1}^x f_k^{xx}}$$

$$\bar{Q}_k^{uu} = l_k^{uu} + (f_k^u)^T V_{k+1}^{xx} f_k^u + \cancel{V_{k+1}^x f_k^{uu}}$$

$$\bar{Q}_k^{ux} = l_k^{ux} + (f_k^u)^T V_{k+1}^{xx} f_k^x + \underbrace{\cancel{V_{k+1}^x f_k^{ux}}}_{\text{second derivative, usually neglected}}$$

second derivative, usually neglected

- the second-order derivatives of the dynamics are often neglected because they take a lot of time to compute (this variant of DDP is sometimes called iLQR or iLQG)

quadratic approximation of the value function

- remember that Q is the argument of the minimization; to find the value function approximation we derive it with respect to Δu_k and set it equal to zero

$$\frac{\partial \tilde{Q}(x_k, u_k)}{\partial \Delta u_k} = \bar{Q}_k^u + \bar{Q}_k^{uu} \Delta u_k + \bar{Q}_k^{ux} \Delta x_k = 0$$

- the optimal Δu_k is therefore

$$\Delta u_k = -(\bar{Q}_k^{uu})^{-1} \left(\bar{Q}_k^u + \bar{Q}_k^{ux} \Delta x_k \right)$$

- this is a **local feedback** control law, which we can write as

$$\Delta u_k = k_k + K_k \Delta x_k \quad \begin{aligned} k_k &= -(\bar{Q}_k^{uu})^{-1} \bar{Q}_k^u \\ K_k &= -(\bar{Q}_k^{uu})^{-1} \bar{Q}_k^{ux} \end{aligned}$$


k_k is a feedforward term, and K_k is a local optimal gain

DDP algorithm

- start with an initial guess trajectory (\bar{x}, \bar{u})
- backward pass:
 - ▶ compute the value function approximation in the terminal state x_N
 - ▶ compute the approximation in the previous state x_{N-1}
 - ▶ in the process, you also obtain the optimal control law (k_k, K_k)
 - ▶ repeat until reaching the start of the horizon
- forward pass:
 - ▶ start from the current state $x_0 = \bar{x}_0$
 - ▶ compute the new input using the local gains $u_k = \bar{u}_k k_k + K_k \Delta x_k$
 - ▶ propagate the dynamics forward $x_{k+1} = f(x_k, u_k)$
 - ▶ repeat until reaching the end of the horizon
- we now have a new guess trajectory $(\bar{x}, \bar{u}) \leftarrow (x, u)$
- repeat backward pass and forward pass until convergence


line search

- the local optimal control law $\Delta u_k = k_k + K_k \Delta x_k$ is only valid for a quadratic approximation, not for the true value function
- in reality, the value function is not quadratic, so this adjustment might overshoot the local minimum, and increase the cost of the new guess instead of decreasing it




line search

- therefore, we modify the forward pass and perform something similar to a **line search** procedure:
 - start with an $\alpha = 1$
 - propagate the dynamics $x_{k+1} = f(x_k, \bar{u}_k + \alpha k_k + K_k \Delta x_k)$
 - evaluate the cost function over the new trajectory
 - if it increased, choose a smaller α and repeat the forward pass
 - if it decreased, accept the new trajectory (line search finished)




computational complexity

- no large matrix factorization is required! DDP scales **linearly** with the number of variables



computational complexity

- compared with the way SQP scales, we have a clear advantage when operating with a long horizon



example

```
import ... [numpy, casadi, model, etc ...]

# parameters
n, m = 4, 1
N = 100
max_ddp_iters = 10
max_line_search_iters = 10
Q = np.eye(n) * 0
R = np.eye(m) * 0.01
Q_ter = np.eye(n) * 10000
x_ter = np.array((math.pi, 0, 0, 0))

# symbolic variables
opt = cs.Opti()
X = opt.variable(n)
U = opt.variable(m)

# ... continues in the next slide →
```

```
# cost function
L_ = lambda x, u: (x_ter - x).T @ Q @ (x_ter - x) + u.T @ R @ u
L_ter_ = lambda x: (x_ter - x).T @ Q_ter @ (x_ter - x)
L      = cs.Function('L'      , [X, U], [L_(X,U)])
L_ter  = cs.Function('L_ter'  , [X]  , [L_ter_(X)])
Lx     = cs.Function('Lx'     , [X, U], [cs.jacobian(L(X,U), X)])
Lu     = cs.Function('Lu'     , [X, U], [cs.jacobian(L(X,U), U)])
Lxx    = cs.Function('Lxx'    , [X, U], [cs.jacobian(Lx(X,U), X)])
Lux    = cs.Function('Lux'    , [X, U], [cs.jacobian(Lu(X,U), X)])
Luu    = cs.Function('Luu'    , [X, U], [cs.jacobian(Lu(X,U), U)])
L_terx = cs.Function('L_terx' , [X]  , [cs.jacobian(L_ter(X), X)])
L_terxx = cs.Function('L_terxx', [X]  , [cs.jacobian(L_terx(X), X)])

# dynamics
f = model.get_pendubot_model()
f = cs.Function('f' , [X, U], [f_(X,U)])
fx = cs.Function('fx', [X, U], [cs.jacobian(f_(X,U), X)])
fu = cs.Function('fu', [X, U], [cs.jacobian(f_(X,U), U)])

# ... continues in the next slide ->
```



```
for iter in range(max_ddp_iters):
    # backward pass
    backward_pass_start_time = time.time()
    V[N] = L_ter(x[:,N])
    Vx[:,N] = L_terx(x[:,N])
    Vxx[:, :, N] = L_terxx(x[:,N])

    for i in reversed(range(N)):
        fx_eval = fx(x[:, i], u[:, i])
        fu_eval = fu(x[:, i], u[:, i])

        Qx = Lx(x[:, i], u[:, i]).T + fx_eval.T @ Vx[:, i+1]
        Qu = Lu(x[:, i], u[:, i]).T + fu_eval.T @ Vx[:, i+1]

        Qxx = Lxx(x[:, i], u[:, i]) + fx_eval.T @ Vxx[:, :, i+1] @ fx_eval
        Quu = Luu(x[:, i], u[:, i]) + fu_eval.T @ Vxx[:, :, i+1] @ fu_eval
        Qux = Lux(x[:, i], u[:, i]) + fu_eval.T @ Vxx[:, :, i+1] @ fx_eval

        Quu_inv = np.linalg.inv(Quu)
        k[i] = - Quu_inv @ Qu
        K[i] = - Quu_inv @ Qux

    V[i] = V[i+1] - 0.5 * k[i].T @ Quu @ k[i]
    Vx[:, i] = np.array(Qx - K[i].T @ Quu @ k[i]).flatten()
    Vxx[:, :, i] = Qxx - K[i].T @ Quu @ K[i]

    # ... continues in the next slide ->
```

```
# forward pass
forward_pass_start_time = time.time()
unew = np.ones((m, N))
xnew = np.zeros((n, N+1))
xnew[:,0] = x[:,0]

# line search
alpha = 1.
for ls_iter in range(max_line_search_iters):
    new_cost = 0
    for i in range(N):
        unew[:,i] = u[:,i] + alpha * k[i] + K[i] @ (xnew[:,i] - x[:,i])
        xnew[:,i+1] = np.array(f(xnew[:,i], unew[:,i])).flatten()
        new_cost = new_cost + L(xnew[:,i], unew[:,i])
    new_cost = new_cost + L_ter(xnew[:,N])

    if new_cost < cost:
        cost = new_cost
        x = xnew
        u = unew
        break
    else:
        alpha /= 2.
```

- **pro**: the computations are propagated forward and backward along the horizon: if you double the horizon length you just double the number of computations
- this means that the computational cost is **linear** with the horizon length
- by contrast, the complexity of SQP depends on the complexity of the QP subproblem, which in turns depends on the solver used (we don't know it apriori)

- **con**: it does not directly allow to set **constraints**, whereas in SQP this is straightforward
- **con**: it is a **single shooting** method, which means that when integrating the initial input guess you could get a diverging trajectory
- **pro**: there are variants of DDP that use an initial guess on the state, incorporating defects in a way that is similar to **multiple shooting** [Mastalli et al., 2020]

Mastalli et al., "Crocodyl: An Efficient and Versatile Framework for Multi-Contact Optimal Control", Robotics and Automation Letters 2020