


# Lecture 7: Humanoid Robot Control

Nicola Scianca

December 2024

# gait phases




- human walking typically consists of the cyclic alternation of **four phases**
- robots typically have flat non-flexible feet, and are usually limited to **single support** and **double support** phases

- to generate the commands necessary to achieve these motions we can use different architectures
- let us take a look at some examples, before exploring in depth a specific architecture
  - ▶ **end-to-end** architectures
  - ▶ architectures that use **whole-body predictive controllers**
  - ▶ architectures that perform predictive control with a **simplified model**

# end-to-end architectures

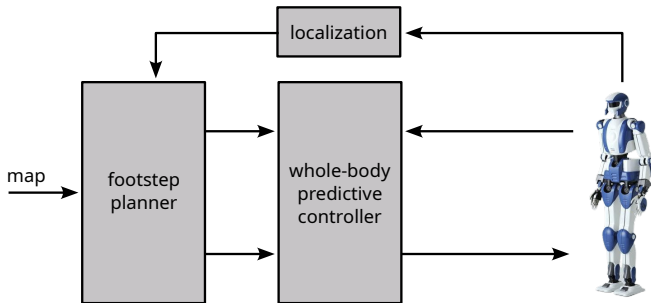
- end-to-end architectures are typically **data-driven** (e.g., based on **reinforcement learning** or **imitation learning**)



- the strength of these techniques lies in the fact that they can directly use **visual information**, e.g., coming from an RGB camera, to perceive the environment
- challenges include the **sim-to-real gap** usually found in learning-based controllers, and the fact that performing a different task usually requires **retraining**
- as of now, these approaches work well on **quadrupeds** but are more challenging on humanoids

# whole-body predictive controllers

- some architectures are based on perform predictive control on the **whole-body model** of the robot

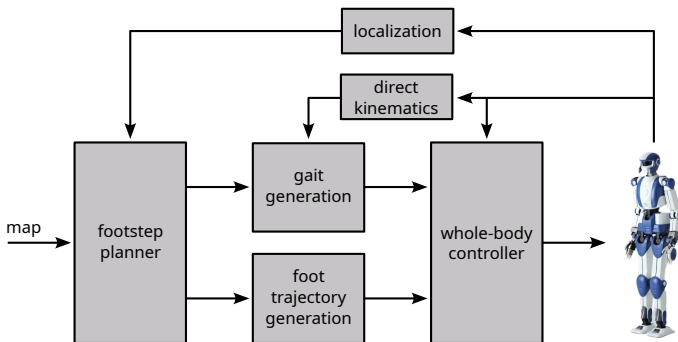


# whole-body predictive controllers

- these techniques are potentially capable of performing very **dynamic** motions, such as running and jumping
- they require **heavy computations**, and often a good **initial guess**, to avoid the optimization getting stuck in a bad **local minimum**
- long-term goals must be planned using a separate technique (e.g., a footstep planner to reach a goal) and some form of localization

# gait generation with a simplified model

- a common approach is to have a separate controller optimize the trajectory using a **simplified model**, that can then be tracked by a whole-body controller





# gait generation with a simplified model


- the range of applicability of these architectures is limited by the range of validity of the assumptions used to derive the **simplified model**
- for achieving locomotion they can be quite effective, and with proper choice of the simplified model they even allow jumping and running
- let us now analyze more in detail a typical architecture of this kind, starting from the whole-body controller

- in this lecture we will focus on **hierarchical** architectures based on **model predictive control** (MPC)
- MPC is used to optimize a trajectory for a **simplified model** of the dynamics
- a whole-body controller is used to generate commands for the robot


# what is a gait?

- a **gait** is the way a person **walks**
- in humanoid robotics, **gait generation** the process by which we generate a trajectory for the **Center of Mass** (CoM) of the robot
- assume we want to go from a start position to a goal position: let us now look at the conceptual steps we have to go through in order to achieve this


- the first thing to do is **plan** a sequence of **footsteps** (e.g., using RRT); this defines the shape and position of the **support polygon** at each time during the walk



- now we can design a **ZMP trajectory** that is always inside the **support polygon**: it uses double support phases to transition between consecutive footsteps




- the next step is to generate a **CoM trajectory** that satisfies the dynamics of the **simplified model** we chose to use (e.g., the LIP model)



# gait generation

- now we generate the other trajectories that we need (e.g., feet, torso orientation, ...) and track them using a **whole-body controller**



- since the ZMP is the input of the LIP model, we could in principle design an arbitrary ZMP trajectory  $p_z(t)$ , plug it inside the model and **integrate** forward to obtain the CoM trajectory  $p_c(t)$

$$\text{ZMP} \implies \boxed{\ddot{p}_c = \eta^2(p_c - p_z)} \implies \text{CoM}$$

- however, the LIP is **unstable!** if we do not keep this into account we would get a **diverging** CoM trajectory
- a solution is to have an **optimization-based controller** design both the **ZMP** and the **CoM** trajectories



- starting from the current time-step  $k$ , we want to predict a **discrete** sequence of input and states over  $N$  future time-steps

$$\mathbf{x}_k, \mathbf{x}_{k+1}, \dots, \mathbf{x}_{k+N-1}, \mathbf{x}_{k+N}$$

$$\mathbf{u}_k, \mathbf{u}_{k+1}, \dots, \mathbf{u}_{k+N-1}$$

- these states and inputs are connected via the systems **dynamics**, which we write in discrete form

$$\mathbf{x}_{i+1} = f(\mathbf{x}_i, \mathbf{u}_i)$$

- for physical systems, which have continuous dynamics, we assume to have discretized them in some way (e.g., Euler, Runge-Kutta, ...)

# model predictive control

- at time-step  $k$ , an **MPC** computes an optimal control sequence by solving an optimization problem of the form

$$\begin{aligned} \min \quad & \sum_{i=k}^{k+N-1} l(\mathbf{x}_i, \mathbf{u}_i) + l_N(\mathbf{x}_N, \mathbf{u}_N) \\ \text{s.t.} \quad & \mathbf{x}_k = \bar{\mathbf{x}}_k \\ & \mathbf{x}_{i+1} = f(\mathbf{x}_i, \mathbf{u}_i) \quad \text{for } i = 0, \dots, N-1 \\ & h(\mathbf{x}_i, \mathbf{u}_i) = 0 \quad \text{for } i = 0, \dots, N \\ & g(\mathbf{x}_i, \mathbf{u}_i) \leq 0 \quad \text{for } i = 0, \dots, N \end{aligned}$$

- the initial state  $\mathbf{x}_k$  of the predicted sequence is set equal to the **measured state**  $\bar{\mathbf{x}}_k$  (**feedback!**)
- other constraints ensure that the **dynamics** are satisfied, and enforce other requirements on states and inputs

- generic optimization problems can be **computationally heavy**, but an MPC controller must give us solutions at a rate of about 100 Hz
- to keep computation time manageable, we typically formulate this as a **Quadratic Program** (QP), i.e., an optimization problem with **quadratic cost** and **linear constraints**
- these problems can be solved using off-the-shelf libraries in **milliseconds**

- MPC gait generation using the LIP model was introduced with the **preview controller** [Kajita et al.]
- this technique used a cost function to track a **reference ZMP** trajectory, designed in such a way to always be inside the support polygon
- it did not make use of **constraints**, therefore the solution was given in **closed form** (without the need to use a QP solver)

Kajita et al., "Biped walking pattern generation by using preview control of zero-moment point", International Conference on Robotics and Automation, 2003

- the dynamics is represented having the **CoM jerk**  $\ddot{p}_c$  (third time derivative) as the input, and the **ZMP** as the output

$$\underbrace{\frac{d}{dt} \begin{pmatrix} p_c \\ \dot{p}_c \\ \ddot{p}_c \end{pmatrix}}_{\dot{x}} = \underbrace{\begin{pmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{pmatrix}}_A \underbrace{\begin{pmatrix} p_c \\ \dot{p}_c \\ \ddot{p}_c \end{pmatrix}}_x + \underbrace{\begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}_B \underbrace{\ddot{p}_c}_u$$

$$\underbrace{p_z}_y = \underbrace{\begin{pmatrix} 1 & 0 & -1/\eta^2 \end{pmatrix}}_C \underbrace{\begin{pmatrix} p_c \\ \dot{p}_c \\ \ddot{p}_c \end{pmatrix}}_x$$

- the jerk  $\ddot{p}_c$  is used instead of the acceleration  $\ddot{p}_c$  to avoid the output directly depending on the input, which might cause discontinuous ZMP trajectories

- the optimization problem can now be written as

$$\begin{aligned} \min_{\mathbf{x}_i, \mathbf{u}_i} \quad & \sum_{i=k}^{k+N-1} \left( \left( \ddot{p}_c^i \right)^2 + \beta \left( p_z^i - p_z^{i,\text{ref}} \right)^2 \right) + \beta \left( p_z^{k+N} - p_z^{k+N,\text{ref}} \right)^2 \\ \text{s.t.} \quad & \mathbf{x}_k = \bar{\mathbf{x}}_k \\ & \mathbf{x}_{i+1} = \mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i \quad \text{for } i = 0, \dots, N-1 \end{aligned}$$

- there are constraints to enforce the consistency of the predicted sequence with the **dynamics** of the model
- $\beta$  is a weight that modulates how strongly we want to **track** the reference

- the predicted states can be expressed in terms of the measured current state  $\bar{x}_k$  and the predicted input sequence  $u_k, \dots, u_{k+N-1}$  by performing **algebraic substitutions**

$$x_{k+1} = A\bar{x}_k + Bu_k$$

$$\begin{aligned} x_{k+2} &= Ax_{k+1} + Bu_{k+1} = A(A\bar{x}_k + Bu_k) + Bu_{k+1} \\ &= A^2\bar{x}_k + ABu_k + Bu_{k+1} \end{aligned}$$

$$\begin{aligned} x_{k+3} &= Ax_{k+2} + Bu_{k+2} = A(A^2\bar{x}_k + ABu_k + Bu_{k+1}) + Bu_{k+2} \\ &= A^3\bar{x}_k + A^2Bu_k + ABu_{k+1} + Bu_{k+2} \end{aligned}$$

$\vdots$

$$x_{k+N} = A^N\bar{x}_k + A^{N-1}Bu_k + \dots + ABu_{k+N-2} + Bu_{k+N-1}$$

- preview control is essentially using **unconstrained MPC** to realize **tracking** of a predefined ZMP trajectory
- the absence of constraints implies that this method is quite fast, because the solution can be recovered in closed form (we have to invert a large matrix but only once)
- however, without constraints we cannot guarantee that the **ZMP** will always be inside the **support polygon**



- since we know the footstep plan, it is possible to enforce **constraints** that will ensure that the ZMP is always contained in the support polygon

$$p_z^{k+i,\min} \leq Cx_{k+i} \leq p_z^{k+i,\max}$$

- if the robot is in **single support**  $p_z^{k+i,\min}$  and  $p_z^{k+i,\max}$  represent the lower and upper edge of the foot
- if the robot is in **double support** the expression is a bit more involved, but still linear

Wieber, "Trajectory free linear model predictive control for stable walking in the presence of strong perturbations", International Conference on Humanoid Robots, 2006

- having constrained the ZMP, we do not need to explicitly track the reference in the cost function (but we still minimize the square of the input to make the problem convex)

$$\min_{\mathbf{x}_i, \mathbf{u}_i} \sum_{i=k}^{k+N-1} \left( \ddot{p}_c^i \right)^2$$

$$\text{s.t.} \quad \mathbf{x}_k = \bar{\mathbf{x}}_k$$

$$\mathbf{x}_{i+1} = \mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i \quad \text{for } i = 0, \dots, N-1$$

$$p_z^{k+i, \min} \leq \mathbf{C}\mathbf{x}_{k+i} \leq p_z^{k+i, \max} \quad \text{for } i = 1, \dots, N$$

- this in principle means that the MPC is free to choose the “best” trajectory, although in practice we might still want to include a ZMP tracking term if it gives better results

- when we derived the LIP model, we saw that it can be written in state-space form by choosing the **ZMP** as the **input**
- to avoid discontinuous ZMP trajectories we can **dynamically extend** the system with the ZMP derivative  $\dot{p}_z$  as new input

$$\underbrace{\frac{d}{dt} \begin{pmatrix} p_c \\ \dot{p}_c \\ p_z \end{pmatrix}}_{\dot{x}} = \underbrace{\begin{pmatrix} 0 & 1 & 0 \\ \eta^2 & 0 & -\eta^2 \\ 0 & 0 & 0 \end{pmatrix}}_A \underbrace{\begin{pmatrix} p_c \\ \dot{p}_c \\ p_z \end{pmatrix}}_x + \underbrace{\begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}}_B \underbrace{\dot{p}_z}_u$$

- since the LIP has a positive eigenvalue, we will get a diverging trajectory if we just set up an MPC that minimizes the square of the input
- to avoid the instability, we can enforce a **constraint** derived from the following **stability condition** on the unstable part of the dynamics [Scianca et al.]

$$\underbrace{p_x^k + \frac{\dot{p}_x^k}{\eta}}_{\text{current state}} = \eta \int_{t_k}^{\infty} e^{-\eta(\tau-t_k)} \underbrace{p_z(\tau)}_{\text{future ZMP}} d\tau$$

- the **future ZMP** trajectory is partly **predicted** by the MPC, and partly **conjectured** based on the footstep plan

Scianca et al., "MPC for Humanoid Gait Generation: Stability and Feasibility", Transactions on Robotics, 2020

- this condition must be enforced inside the optimization problem as an additional **stability constraint**

$$\min_{\mathbf{x}_i, \mathbf{u}_i} \sum_{i=k}^{k+N-1} \left( \dot{p}_z^i \right)^2$$

$$\text{s.t.} \quad \mathbf{x}_k = \bar{\mathbf{x}}_k$$

$$\mathbf{x}_{i+1} = \mathbf{A}\mathbf{x}_i + \mathbf{B}\mathbf{u}_i \quad \text{for } i = 0, \dots, N-1$$

ZMP constraints

stability constraint

- however, we can find guarantees of **feasibility** and **stability** thanks to this additional constraint

# MPC locomotion example

- a working implementation of this last MPC scheme is available at <https://github.com/DIAG-Robotics-Lab/ismpc>
- it is written in Python using the DART robotics toolkit
- the optimization problems are solved with OSQP, interfaced using CasADi

# MPC locomotion example

