

Underactuated Robots

Lecture 2: Trajectory Optimization

Nicola Scianca

December 2024

why optimization

- with complex systems, such as underactuated robots with many degrees of freedom, we often can't explicitly write a control law $u = f(x)$ as a closed-form expression
- what we can do instead is to write an **optimization problem** that encodes our control objectives, and its solution will implicitly define our control law
- these lectures will be mostly focused on optimization-based techniques, which we collectively call **trajectory optimization**

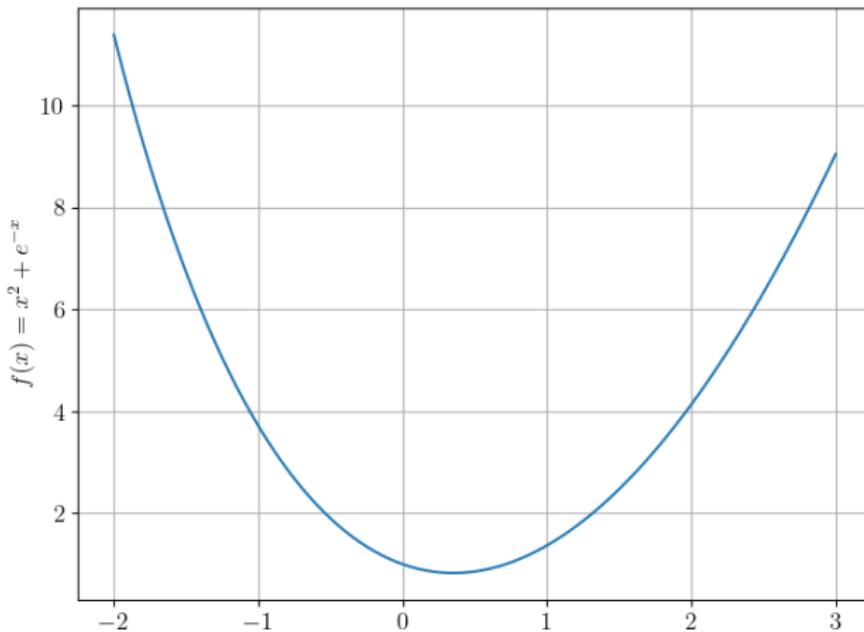
introduction to optimization

- optimization is the process of finding the **minimum** (or the maximum) of some **function**, possibly in the presence of **constraints**
- key concepts:
 - ▶ constrained / unconstrained optimization
 - ▶ convexity
 - ▶ gradient descent
 - ▶ newton's method

unconstrained optimization

- let's say we want to find the minimum of the function

$$f(x) = x^2 + e^{-x}$$



gradient descent

- we learned in calculus class that we can compute the derivative and set it equal to zero

$$\frac{df}{dx} = 2x - e^{-x} = 0 \implies 2x = e^{-x}$$

but solving this equation is not straightforward: we must zero it **numerically**

- the basic idea of **gradient descent**: pick some point and check the derivative
 - ▶ if **positive**, the function decreases to the left
 - ▶ if **negative**, the function decreases to the right
- if we just keep going in the direction in which the function **decreases**, we should, at some point, find a (local) minimum

gradient descent

- in gradient descent we are computing a **linear** approximation of the function around our guess \bar{x}

$$f(\bar{x} + \Delta x) = f(\bar{x}) + \nabla f|_{\bar{x}} \cdot \Delta x$$

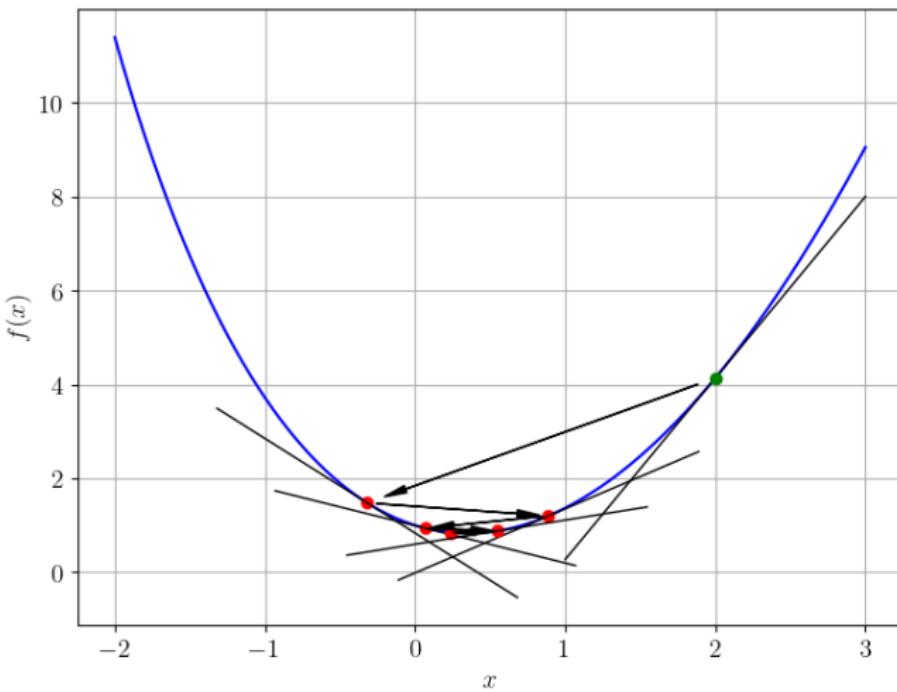
- then we take a step in the direction in which this linear approximation decreases, which means, the negative of its gradient

$$\Delta x = -\alpha \nabla f|_{\bar{x}}$$

- α is the **step size**

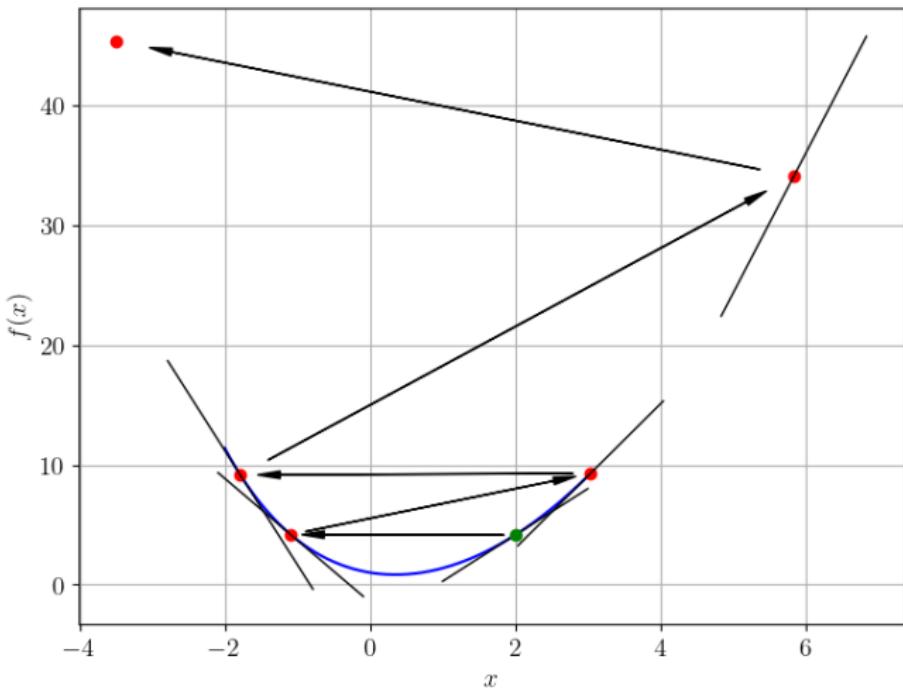
gradient descent

- for our example function, if we start from the point $x = 2$ and iterate with step size $\alpha = 0.6$, we converge quickly



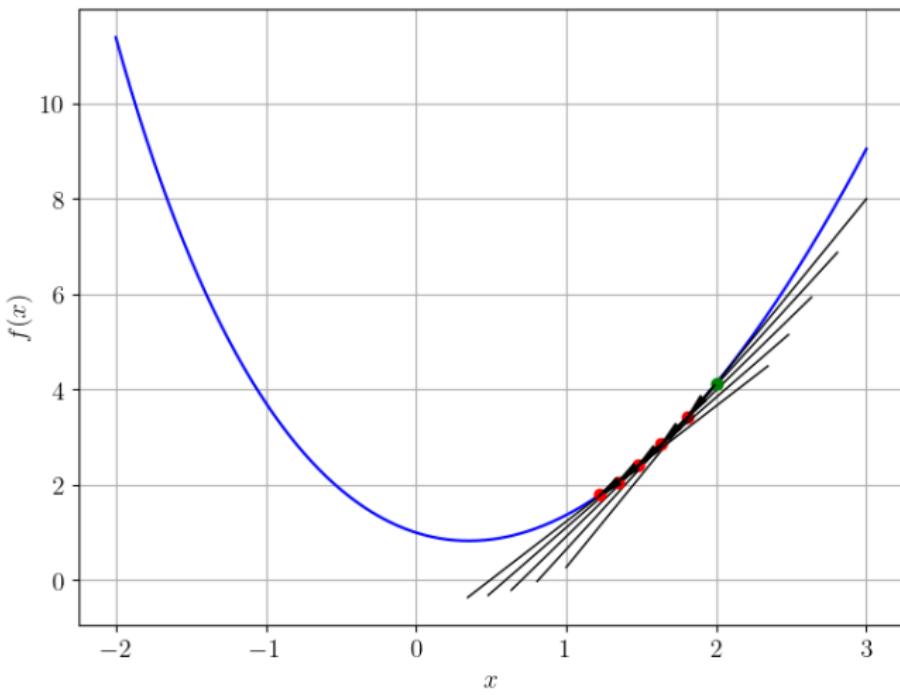
gradient descent

- with just a slightly larger step size ($\alpha = 0.8$) we actually diverge from the minimum!



gradient descent

- if we play it safe and choose a small step size ($\alpha = 0.05$) then the convergence can be very slow!



line search

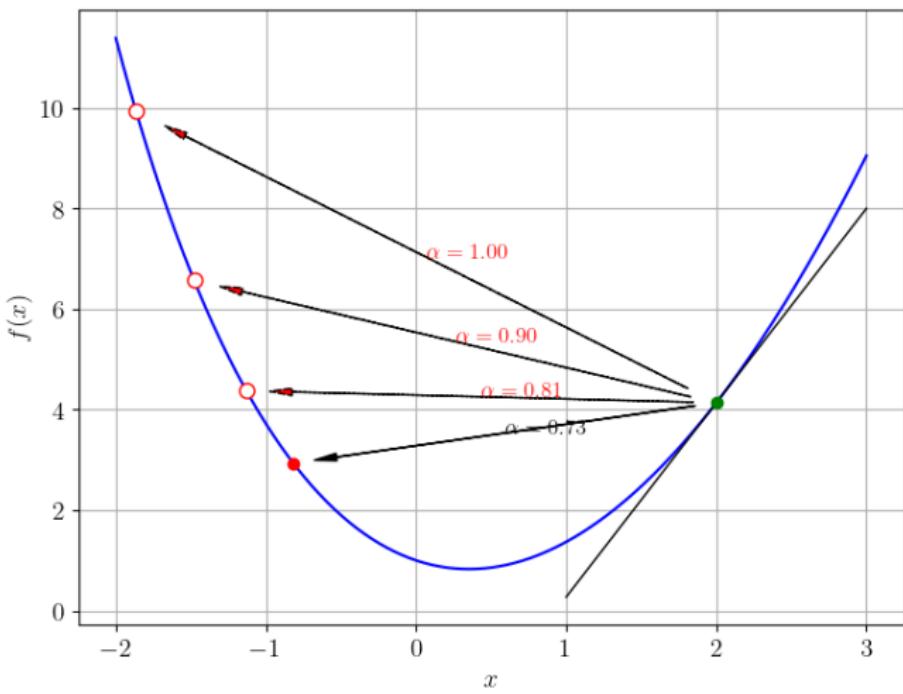
- one way to avoid these problems is to do a **line search**
- the basic idea is that we want to move **along the descent direction** (the negative gradient), but we want to choose the step size small enough to have the function decrease
- the simplest way: starting from an initial step size, say $\alpha = 1$ we check if our function decreases with this step size

$$f(\bar{x} - \alpha \nabla f|_{\bar{x}}) \leq f(\bar{x})$$

- if it does, then we **accept** the step, if it doesn't, we reduce α and try again

gradient descent

- in this example we reduce α by a factor of 0.9 each time
(usually it is reduced by 0.5, this is just for illustration)



Newton's method

- Newton's method is another technique for numerically finding the minimum of a function
- instead of just looking at the descent direction, at each iteration we approximate the function $f(x)$ as a **quadratic function**, then jump to the minimum of the approximation
- it uses more information (the second derivative), but usually performs much better

Newton's method

- first we have to compute both the first and second derivative

$$\frac{df}{dx} = 2x - e^{-x}, \quad \frac{d^2f}{dx^2} = 2 + e^{-x}$$

- this gives us a **quadratic** approximation of our function (second order Taylor expansion)

$$f(\bar{x} + \Delta x) = f(\bar{x}) + \left. \frac{df}{dx} \right|_{\bar{x}} \Delta x + \frac{1}{2} \left. \frac{d^2f}{dx^2} \right|_{\bar{x}} \Delta x^2$$

- the minimum of this approximation is easy to find

$$\left. \frac{df}{dx} \right|_{\bar{x}} + \left. \frac{d^2f}{dx^2} \right|_{\bar{x}} \Delta x = 0 \quad \Rightarrow \quad \Delta x = - \left(\left. \frac{d^2f}{dx^2} \right|_{\bar{x}} \right)^{-1} \left. \frac{df}{dx} \right|_{\bar{x}}$$

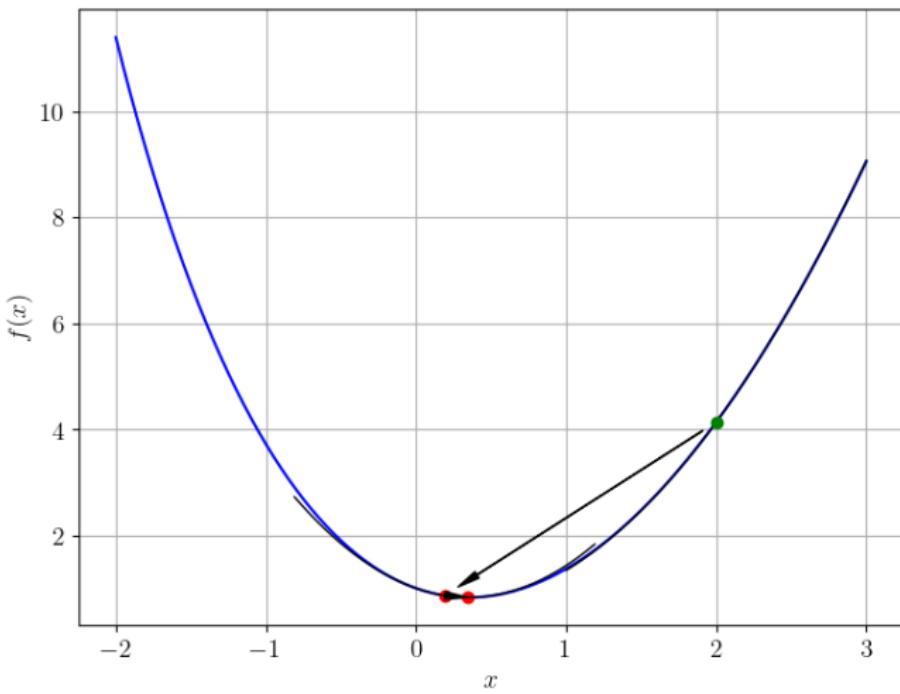
Newton's method

$$\Delta x = - (\nabla^2 f|_{\bar{x}})^{-1} \nabla f|_{\bar{x}}$$

- note that, if $(\nabla^2 f|_{\bar{x}})^{-1}$ is positive, Newton's method also follows the direction of the negative gradient, but it does so in a smarter way
- however, if $(\nabla^2 f|_{\bar{x}})^{-1}$ is negative, we actually go up the gradient, and we might end up in a maximum!
- this should make sense because it means that our quadratic approximation is upside down, so Newton's method actually takes us to the maximum

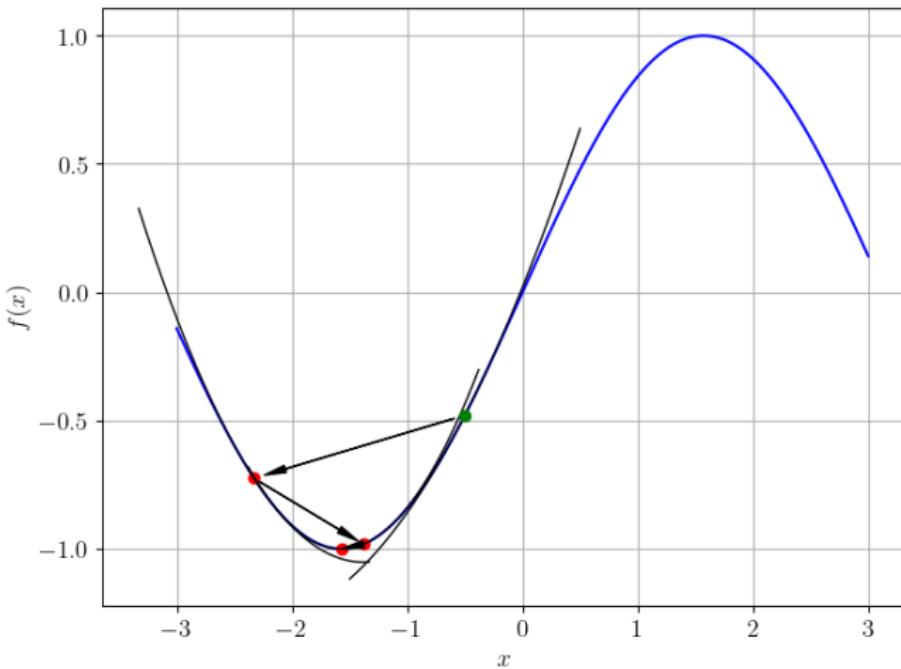
Newton's method

- in our example, Newton's method converges very fast, and we don't even have to choose a step size!



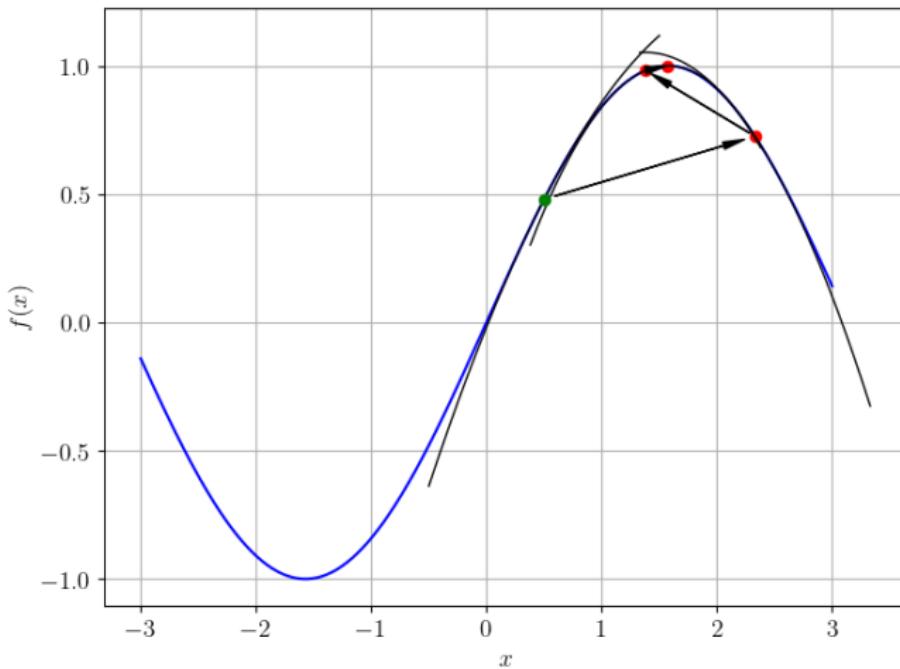
Newton's method

- however, look at what happens if we have a **non-convex** function: if we start close to the minimum, everything is fine



Newton's method

- but if we start at a point where the function is not **convex**, we actually end up in a maximum!



convexity

- **convexity**: if we pick any two points x_a and x_b , and draw a line in between them, the function $f(x)$ is **below** this line

$$\forall x_a, x_b \in \mathbb{R}, \alpha \in (0, 1)$$

$$f(\alpha x_a + (1 - \alpha)x_b) < \alpha f(x_a) + (1 - \alpha)f(x_b)$$

- convex functions have a **single minimum**, thus if we find a solution we can always be sure that it's the **global optimum**
- also, convex functions have a Hessian $\nabla^2 f$ that is always positive, which means that we never run into trouble with Newton's method (if not: regularization, Gauss-Newton approximation, etc...)

trajectory optimization

- **trajectory optimization** is a mathematical approach to compute the best trajectory for a system to follow, while satisfying **constraints** and minimizing a **cost function**
- it involves optimizing the states and control inputs over time to achieve a specific goal, such as minimizing energy, time, effort, or following a desired path
- trajectory optimization can be used, in many forms, for tasks like motion planning, navigation, and control

trajectory optimization: examples

- controlling a **fixed-base manipulator** to minimize input torques while avoiding workspace obstacles



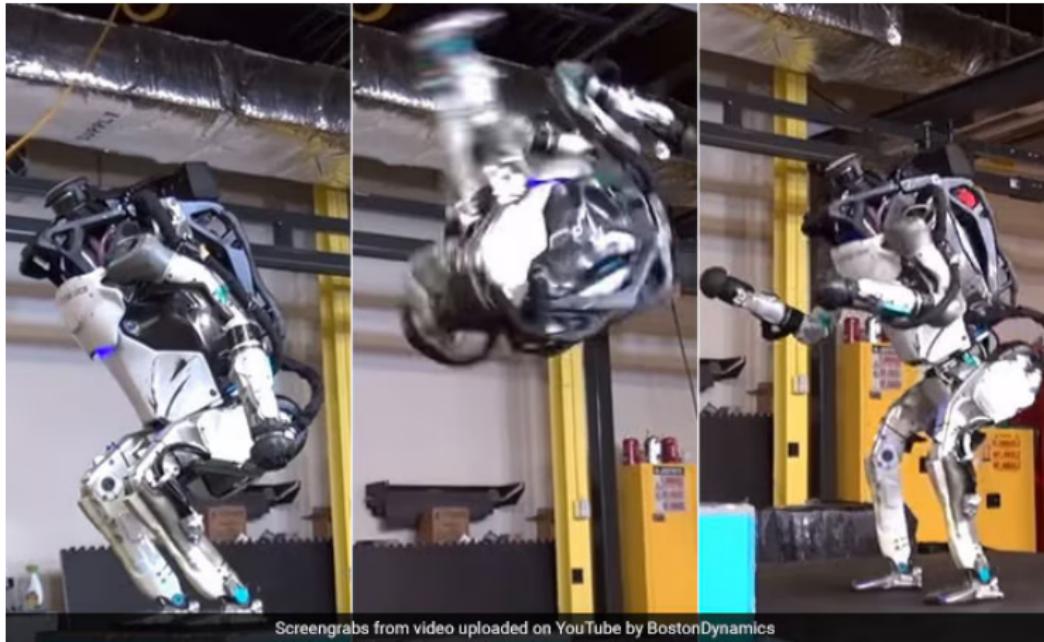
- controlling an **autonomous vehicle** to minimize the distance from the center of the road while avoiding incidents



- controlling a **legged robot** to follow a given reference velocity and making sure it doesn't fall



trajectory optimization: examples



Screengrabs from video uploaded on YouTube by BostonDynamics

<https://www.youtube.com/embed/fRj34o4hN4I>

trajectory optimization

- in the continuous case, we can formulate a basic trajectory optimization problem as

$$\min \int_{t_0}^{t_f} l(x(t), u(t), t) dt + l_t(x(t_f))$$

subject to

$$\dot{x} = f(x(t), u(t)), \quad \text{for } t \in [t_0, t_f],$$

$$x(t_0) = \bar{x}_0,$$

$$h(x(t), u(t)) = 0, \quad \text{for } t \in [t_0, t_f],$$

$$g(x(t), u(t)) \leq 0, \quad \text{for } t \in [t_0, t_f],$$

trajectory optimization

- in this problem we have a **cost function** composed of a **running cost** integrated along the trajectory, and a **terminal cost** on the final state
- we also have constraints on the **initial state** and on the **system dynamics**
- we can also add other **equality constraints** $h(x(t), u(t)) = 0$ and **inequality constraints** $g(x(t), u(t)) \leq 0$, for example to perform obstacle avoidance, or to limit control effort

system dynamics discretization

- we consider a **discrete** sequence of input and states over N time-steps, covering the entire duration of the task

$$x_0, x_1, \dots, x_{N-1}, x_N$$

$$u_0, u_1, \dots, u_{N-1}$$

- these states and inputs are connected via the systems **dynamics**, which we write in discrete form

$$x_{k+1} = f(x_k, u_k)$$

- for physical systems, which have continuous dynamics, we assume to have discretized them in some way (e.g., Euler, Runge-Kutta, ...)

cost function

- to formulate a trajectory optimization problem, we choose a **cost function**, which is something that we want to minimize

$$\min \sum_{i=0}^{N-1} \underbrace{l_k(x_i, u_i)}_{\text{running cost}} + \underbrace{l_N(x_N)}_{\text{terminal cost}}$$

- for example, the cost function might be a measure of energy consumption, or the time necessary to complete the task, or the distance to some reference point or trajectory
- the cost function is usually constituted by a **running cost**, evaluating the cost of states and inputs along the trajectory, and a **terminal cost** evaluating the cost of the final state

constraints

- this cost function must be minimized while taking into account some necessary **constraints**

$$x_{i+1} = f(x_i, u_i) \quad \text{for } i = 0, \dots, N - 1$$

$$x_0 = \bar{x}_0$$

- the input and state sequences should satisfy the dynamics, which makes it so that the resulting input and state sequences constitute legitimate **trajectories** for the system
- the initial state should correspond to the initial state of the system \bar{x}_0
- we could add additional constraints (e.g., max input, obstacle avoidance, ...)

trajectory optimization

- a basic trajectory optimization problem has the following form

$$\min \sum_{i=0}^{N-1} l_k(x_i, u_i) + l_N(x_N)$$

$$\begin{aligned} \text{s. t. } \quad & x_{i+1} = f(x_i, u_i) \quad \text{for } i = 0, \dots, N-1 \\ & x_0 = \bar{x}_0 \end{aligned}$$

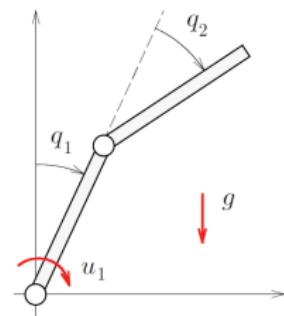
- the optimal input that we get by solving this problem depends on the current state \bar{x}_0
- if we could solve this problem for **every** state, we would have found an **optimal control law** $u(x)$

trajectory optimization

- if we can't explicitly write out a feedback law $u(x)$, maybe we can still solve the optimization problem, starting from one initial state \bar{x}_0 , and find a particular trajectory as solution
- by doing this we don't get an optimal control law, but rather a particular **optimal trajectory**, i.e., a sequence of inputs and states which satisfy the system dynamics
- we can perform this optimization **offline**, and then do trajectory tracking to keep the system as close as possible to this optimal trajectory

example

- consider a **pendubot**: a double pendulum with one actuated joint
- the state variables are
 - ▶ θ_1 angle of the first link
 - ▶ θ_2 angle of the second link (relative)
 - ▶ $\dot{\theta}_1$ angular velocity of the first link
 - ▶ $\dot{\theta}_2$ angular velocity of the second link
- goal: swing up the pendubot to its upright configuration $(0, 0, 0, 0)$
- initial condition: stable equilibrium $(\pi, 0, 0, 0)$ in which both arms of the pendulum are pointing down



example

- first, we discretize states and inputs

$$u_0, u_1, \dots, u_{N-1}$$

$$x_0, x_1, \dots, x_{N-1}, x_N$$

- then, we define a cost function, which in our case is just to minimize the sum of the squares of the input

$$\min \sum_{i=0}^{N-1} u_i^2$$

- and we define our initial state (given) and our final state (target) as constraints

$$x_0 = (\pi, 0, 0, 0)$$

$$x_N = (0, 0, 0, 0)$$

example

- the dynamics of the pendubot can be written as

$$M(q)\ddot{q} + n(q, \dot{q}) = \underbrace{(u, 0)^T}_{\text{zero component due to the underactuation}}$$

- first we rewrite it as $\dot{x} = f(x, u)$, with $x = (q, \dot{q})$

$$\dot{x} = \left(M(q)^{-1} \left(\begin{pmatrix} \dot{q} \\ u \end{pmatrix} - n(q, \dot{q}) \right) \right) = f(x, u)$$

- we must discretize the dynamics, e.g. with Euler integration

$$x_{k+1} = x_k + \Delta t \cdot f(x_k, u_k)$$

transcription methods

- **transcription** is the particular way in which we choose we write our optimization problem so that we can then solve it
- there are many transcription methods: shooting methods, collocation methods, ...
- we will discuss the two main paradigms that characterize shooting methods: **single shooting** and **multiple shooting**
- aside from very simple cases, the solution process will always involve **numerical iteration** that hopefully **converge** to an optimal solution

transcription methods

- the same problem can be **transcribed** in different ways: we will consider the two main possibilities (but there are others)
- **single shooting** methods: we perform substitutions to have less variables and constraints, but we increase the nonlinearity
- **multiple shooting** methods: we have a larger number of variables held together by constraints, so the problem is larger but less nonlinear

transcription methods: single shooting

- in **single shooting** methods, the **inputs** u_0, \dots, u_{N-1} are decision variables: we eliminate states from the cost function by substituting in the dynamics

$$x_1 = f(\bar{x}_0, u_0)$$

$$x_2 = f(x_1, u_1) = f(f(\bar{x}_0, u_0), u_1)$$

...

$$x_k = f(f(\dots f(\bar{x}_0, u_0), u_1), u_2), \dots, u_{k-1})$$

- we also substitute x_0 with \bar{x}_0 to eliminate the initial state constraint
- we get a problem that is only dependent on the input sequence, and has **no constraints**

transcription methods: multiple shooting

- in **multiple shooting** methods, the **inputs** and **states** are decision variables: they are linked together by constraints (the dynamic equation)

$$x_1 = f(x_0, u_0)$$

$$x_2 = f(x_1, u_1)$$

...

$$x_k = f(x_{k-1}, u_{k-1})$$

- now we have to solve a constrained problem, but the advantage is that we have a simpler cost function
- if we stop the solution before convergence, the dynamics might not be satisfied exactly

transcription methods: example

consider this example with horizon length $N = 2$

- discrete inputs and states

$$u_0, u_1,$$

$$x_0, x_1, x_2$$

- cost function

$$\sum_{i=0}^1 \left(u_i^2 + x_i^2 \right) + x_2^2$$

- dynamics

$$x_{k+1} = f(x_k, u_k)$$

transcription methods: example

- **single shooting** formulation

$$\min_{u_0, u_1} \left(u_0^2 + u_1^2 + \bar{x}_0^2 + f(\bar{x}_0, u_0)^2 + f(f(\bar{x}_0, u_0), u_1)^2 \right)$$

- **multiple shooting** formulation

$$\min_{\substack{u_0, u_1, \\ x_0, x_1, x_2}} \left(u_0^2 + u_1^2 + x_0^2 + x_1^2 + x_2^2 \right)$$

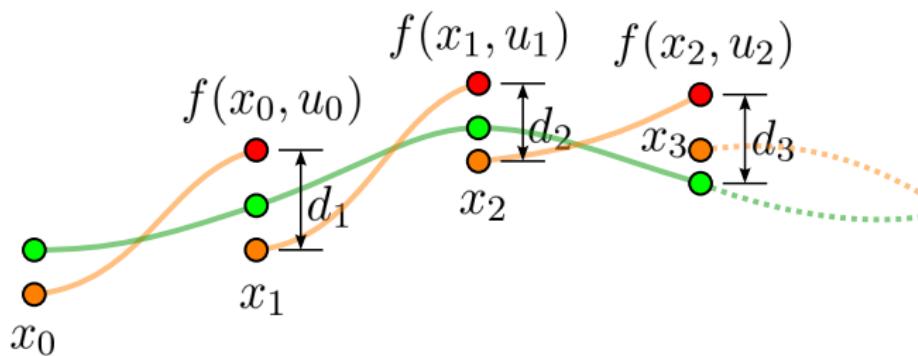
$$\text{s. t. } x_0 = \bar{x}_0$$

$$x_1 = f(x_0, u_0)$$

$$x_2 = f(x_1, u_1)$$

multiple shooting: defects

- in a multiple shooting formulation, the constraints are not necessarily satisfied until we reach convergence
- the difference between the forward integration and the next variable is called the **defect** $d_{k+1} = x_{k+1} - f(x_k, u_k)$



single shooting: numerical problems

- in single shooting, the entire trajectory is determined by forward integration from an initial guess for the control inputs
- this can lead to **numerical instability**, especially for long time horizons or highly nonlinear systems
- the initial trajectory might be too far away from the optimum, or in the worst case we might even have a software failure (numbers get too big)

single shooting vs. multiple shooting

- let's recap the main differences between single and multiple shooting methods

single shooting	multiple shooting
less variables	more variables
cost function is very nonlinear	cost function is usually quadratic
less constraints	more constraints
initial guess could diverge	initial guess is stable
result is always a feasible trajectory	result could have defects

- if possible, it is usually better to do multiple shooting rather than single shooting